

MariaDB Komplettkurs

Agenda

1. Architecture of MariaDB

- [Architecture Server \(Steps\)](#)
- [Query Cache Usage and Performance](#)
- [Optimizer-Basics](#)
- [Storage Engines](#)

2. Installation

- [Installation \(Ubuntu\)](#)
- [Start/Status/Stop von MariaDB](#)
- [Does mariadb listen to the outside world](#)

3. Configuration

- [Adjust configuration and restart](#)
- [Set global server system variable](#)

4. Information Schema / Status / Processes

- [Show server status](#)
- [Kill long running process](#)
- [Kill \(kickout user\) and stop server](#)

5. Security and User Rights

- [Get Rights of user](#)
- [Secure with SSL server/client](#)
- [Create User/Grant/Revoke - Management of users](#)
- Authentication and Authorization
- [User- and Permission-concepts \(best-practice\)](#)

6. Database - Objects

- [Create Database](#)
- [Show all tables within db](#)

7. InnoDB - Storage Engine

- [InnoDB - Storage Engine - Structure](#)
- [Important InnoDB - configuration - options to optimized performance](#)

8. Training Data

- [Setup training data "contributions"](#)

9. Performance

- [io-Last/CPU-Last](#)
- [Views and performance](#)
- [Partitions and Explain](#)
- [3 Phases of DataSize](#)

10. Backup and Restore (Point-In-Time aka PIT)

- [Backup with mysqldump - best practices](#)

- [Flashback](#)
- [mariabackup](#)
- [Use xtrabackup for MariaDB 5.5](#)

11. Optimal use of indexes

- Index-Types
 - [Describe and indexes](#)
 - [Find out indexes](#)
- [Index and Functions \(Cool new feature in MySQL 5.7\)](#)
- [Index and Likes](#)
- [profiling-get-time-for-execution-of.query](#)
- [Find out cardinality without index](#)

12. Monitoring

- [What to monitor?](#)

13. Replication

- [Slave einrichten -gtid](#)
- [Slave einrichten - master_pos](#)
- [MaxScale installieren](#)
- [Reference: MaxScale-Proxy mit Monitoring](#)
- [Walkthrough:Automatic Failover Master Slave](#)

14. Tools

- [Percona-toolkit-Installation](#)
- [pt-query-digest - analyze slow logs](#)
- [pt-online-schema-change howto](#)

15. Extras

- [User Variables](#)
- [Installation sakila-db](#)

16. Documentation

- [Server System Variables](#)
- [MySQL - Performance - PDF](#)
- [Source-Code MariaDB](#)

17. Diagnosis and measurement of performance

- [Best practices to narrow down performance problems](#)

18. Performance and optimization of SQL statements

- [Do not use '*' whenever possible](#)
- [Be aware of subselects - Example 1](#)
- [Optimizer-hints \(and why you should not use them\)](#)

19. Replication

- [Replikation Read/Write](#)

20. Performance

- [Best Practices](#)
- [Example sys-schema and Reference](#)

- [Change schema online \(pt-online-schema-change\)](#)
- [Optimizer-Hints](#)

21. Documentation / Literature

- [Effective MySQL](#)
- [Last Training](#)
- [MySQL - Performance - PDF](#)
- [MariaDB Galera Cluster](#)
- [MySQL Galera Cluster](#)

22. Questions and Answers

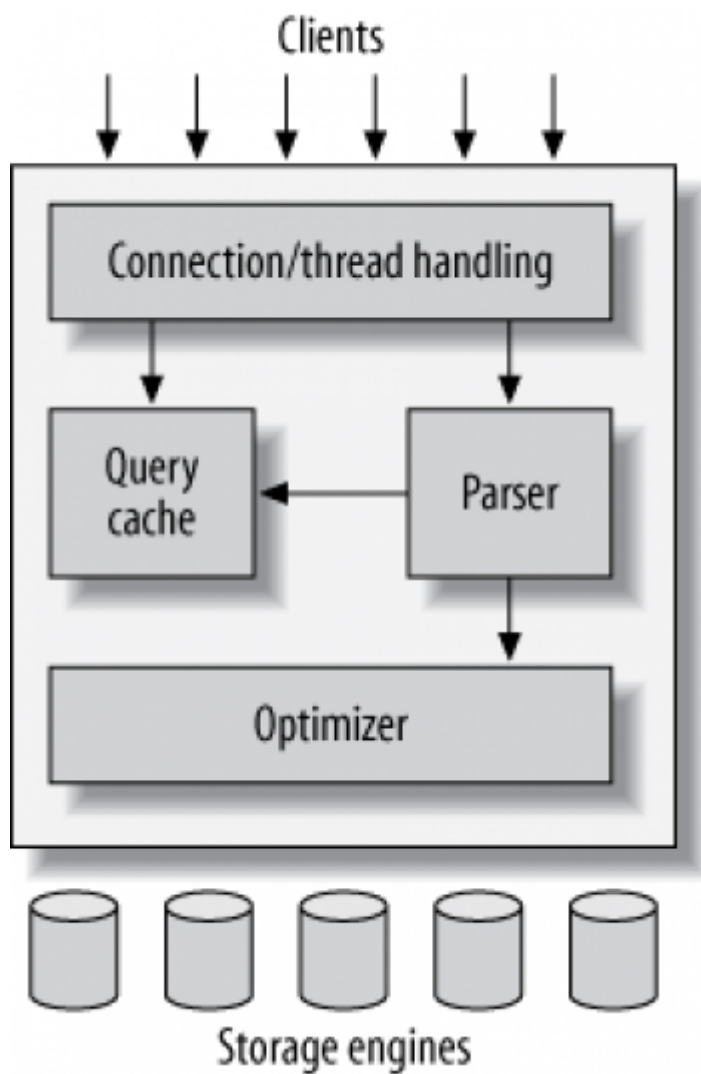
- [Questions and Answers](#)
- [migration-mysql-update-5.6->5.7](#)

23. MySQL Do-Nots

- [mysql-do-nots](#)

Architecture of MariaDB

Architecture Server (Steps)



Query Cache Usage and Performance

Performance query cache

- Always try to optimize innodb with disabled query cache first (innodb_buffer_pool)
- If you use query_cache system can only use on CPU-Core. !!

How to enable query cache

```
## have_query_cache means compiled in mysql
## query_cache_type off means not enable by config
-- query cache is disabled
mysql> show variables like '%query_cache%';
+-----+-----+
| Variable_name          | Value    |
+-----+-----+
| have_query_cache       | YES      |
| query_cache_limit      | 1048576  |
| query_cache_min_res_unit | 4096     |
| query_cache_size       | 1048576  |
| query_cache_type       | OFF      |
| query_cache_wlock_invalidate | OFF      |
+-----+-----+
6 rows in set (0.01 sec)
```

```
root@trn01:/etc/mysql/mysql.conf.d# tail mysqld.cnf
[mysqld]
pid-file           = /var/run/mysqld/mysqld.pid
socket             = /var/run/mysqld/mysqld.sock
datadir            = /var/lib/mysql
log-error          = /var/log/mysql/error.log
## By default we only accept connections from localhost
bind-address       = 0.0.0.0
## Disabling symbolic-links is recommended to prevent assorted security risks
symbolic-links=0
query-cache-type=1

systemctl restart mysql
```

```
mysql> show variables like '%query_cache%';
+-----+-----+
| Variable_name          | Value    |
+-----+-----+
| have_query_cache       | YES      |
| query_cache_limit      | 1048576  |
| query_cache_min_res_unit | 4096     |
| query_cache_size       | 1048576  |
| query_cache_type       | ON       |
| query_cache_wlock_invalidate | OFF      |
+-----+-----+
6 rows in set (0.01 sec)
```

```
mysql> show status like '%Qcache%';
+-----+
| Variable_name          | Value          |
+-----+
| Qcache_free_blocks     | 1              |
| Qcache_free_memory     | 1031832        |
| Qcache_hits            | 0              |
| Qcache_inserts         | 0              |
| Qcache_lowmem_prunes   | 0              |
| Qcache_not_cached      | 0              |
| Qcache_queries_in_cache | 0              |
| Qcache_total_blocks    | 1              |
+-----+
8 rows in set (0.00 sec)

## status in session zurücksetzen.
mysql> flush status;
Query OK, 0 rows affected (0.00 sec)
```

Performance bottleneck - mutex

<https://mariadb.com/de/resources/blog/flexible-mariadb-server-query-cache/>

Something planned ?

- Nope ;o(Demand is new
- You might be able to use Demand together with maxscale
- Refer to: <https://mariadb.com/de/resources/blog/flexible-mariadb-server-query-cache/>

A mutual exclusion object (mutex) is a programming object that allows multiple program threads to share a resource (such as a folder) but not simultaneously. Mutex is set to unlock when the data is no longer needed or when a routine is finished. Mutex creates a bottleneck effect. The blocking means only one query can look at the Query Cache at a time and other queries must wait. A query that must wait to look in the cache only to find it isn't in the cache will be slowed instead of being accelerated.

Optimizer-Basics

General

- All optimizer today are cost-based

Cost-Based

```
## How much costs are needed to get the information
```

Storage Engines

Why ?

```
Let's you choose:  
How your data is stored
```

What ?

- Performance, features and other characteristics you want

What do they do ?

- In charge for: Responsible for storing and retrieving all data stored in MySQL
- Each storage engine has its:
 - Drawbacks and benefits
- Server communicates with them through the storage engine API
 - this interface hides differences
 - makes them largely transparent at query layer
 - api contains a couple of dozen low-level functions e.g. "begin a transaction", "fetch the row that has this primary key"

Storage Engine do not

- Storage Engines do not parse SQL
- Storage Engines do not communicate with each other

They simply

- They simply respond to requests from the server

Which are the most important one ?

- MyISAM/Aria
- InnoDB
- Memory
- CSV
- Blackhole (/dev/null)
- Archive
- Partition
- Federated/FederatedX

Installation

Installation (Ubuntu)

- <https://downloads.mariadb.org/mariadb/repositories/>

```
### repo
sudo apt-get install software-properties-common
sudo apt-key adv --fetch-keys 'https://mariadb.org/mariadb_release_signing_key.asc'
## does an apt update after setting repo - automatically
sudo add-apt-repository 'deb [arch=amd64,arm64,ppc64el]
https://mirror.dogado.de/mariadb/repo/10.5/ubuntu focal main'
sudo apt install mariadb-server
```

Secure installation

```
mariadb-secure-installation
## OR: if not present before 10.4
mysql_secure_installation
```

Start/Status/Stop von MariaDB

```
## als root - user
systemctl status mariadb
systemctl stop mariadb
systemctl start mariadb

<div class="page-break"></div>

### Does mariadb listen to the outside world
```

ls -l | grep mariadb

localhost means it does NOT listen to the outside now

mariabdb 5208 mysql 19u IPv4 56942 0t0 TCP localhost:mysql (LISTEN)

```
<div class="page-break"></div>

## Configuration

### Adjust configuration and restart
```

change config in /etc/mysql/50-server.cnf

After that restart server - so that it takes your new config

systemctl restart mariadb echo \$? # Was call restart succesful -> 0

```
<div class="page-break"></div>

### Set global server system variable

### Find out current value
```

show global variable

show global variables like '%automatic_sp%'

or // variable_name needs to be in capitals

use information_schema select * from global_variables where variable_name like '%AUTOMATIC_SP%';

If you know the exact name

```
select @@global.automatic_sp_privileges; select @@GLOBAL.automatic_sp_privileges;
```

```
### Set global Variable
```

will be set like so till next restart of mysql server

```
set global automatic_sp_privileges = 0
```

```
### automatic_sp_privileges can only be set globally
```

Refer to: server system variable doku

Has same value in global an session scope

```
MariaDB [information_schema]> select @@automatic_sp_privileges; select
@@global.automatic_sp_privileges; +-----+ | @@automatic_sp_privileges | +-----
-----+ | 0 | +-----+ 1 row in set (0.000 sec)

+-----+ | @@global.automatic_sp_privileges | +-----
---+ | 0 | +-----+ 1 row in set (0.000 sec)
```

```
### Reference:
```

```
* https://mariadb.com/kb/en/server-system-variables/#automatic\_sp\_privileges
```

```
<div class="page-break"></div>
```

```
## Information Schema / Status / Processes
```

```
### Show server status
```

```
### Through mysql
```

in mysql interface (client)

```
mysql status;
```

```
### With mysqladmin
```

```
mysqladmin status
```

or if you want to know more

```
mysqladmin extended status
```

```
<div class="page-break"></div>
```

```
### Kill long running process
```

Session 1

sleep for 120 seconds

```
select sleep(120)
```

Session 2

```
show processlist
```

kill process you have identified for sleep(120)

```
MariaDB [(none)]> show processlist; +----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+ | Id | User | Host | db | Command | Time | State | Info | Progress | +----+----
+-----+-----+-----+-----+-----+-----+-----+-----+ | 36 | root |
localhost | NULL | Query | 0 | starting | show processlist | 0.000 | | 37 | root | localhost | training | Query | 4
| User sleep | select sleep(120) | 0.000 | +----+-----+-----+-----+-----+-----+-----+
+-----+-----+ 2 rows in set (0.000 sec)
```

take 37

```
kill 37
```

Session 1: query terminates

ERROR 2013 (HY000): Lost connection to MySQL server during query

```
<div class="page-break"></div>
```

```
### Kill (kickout user) and stop server
```

```
MariaDB [mysql]> show processlist; +----+-----+-----+-----+-----+-----+-----+-----
+-----+-----+ | Id | User | Host | db | Command | Time | State | Info | Progress | +----+----
+-----+-----+-----+-----+-----+-----+-----+-----+ | 30 | root |
localhost | mysql | Sleep | 10 | | NULL | 0.000 | | 34 | root | localhost | mysql | Query | 0 | starting | show
processlist | 0.000 | | 43 | training | localhost | training | Sleep | 5 | | NULL | 0.000 | +----+-----+-----
+-----+-----+ 3 rows in set (0.000 sec)
```

```
MariaDB [mysql]> quit Bye root@its-lu20s04:~# mysql -e 'kill 43' && systemctl stop mariadb root@its-
lu20s04:~#
```

```
<div class="page-break"></div>
```

```
## Security and User Rights

### Get Rights of user

### Secure with SSL server/client

### Create CA and Server-Key
```

On Server - create ca and certificates

```
sudo mkdir -p /etc/mysql/ssl sudo cd /etc/mysql/ssl
```

create ca.

```
sudo openssl genrsa 4096 > ca-key.pem
```

create ca-certificate

Common Name: MariaDB Admin

```
sudo openssl req -new -x509 -nodes -days 365000 -key ca-key.pem -out ca-cert.pem
```

create server-cert

Common Name: MariaDB Server

Password: --- leave empty ----

```
sudo openssl req -newkey rsa:2048 -days 365000 -nodes -keyout server-key.pem -out server-req.pem
```

Next process the rsa - key

```
sudo openssl rsa -in server-key.pem -out server-key.pem
```

Now sign the key

```
sudo openssl x509 -req -in server-req.pem -days 365000 -CA ca-cert.pem -CAkey ca-key.pem -set_serial 01 -out server-cert.pem
```

```
### Verify certificates
```

```
openssl verify -CAfile ca-cert.pem server-cert.pem
```

```
### Configure Server
```

create file

/etc/my.cnf.d/z_ssl.cnf

```
[mysqld] ssl-ca=/etc/mysql/ssl/ca-cert.pem ssl-cert=/etc/mysql/ssl/server-cert.pem ssl-key=/etc/mysql/ssl/server-key.pem
```

Set up TLS version here. For example TLS version 1.2 and 1.3

```
tls_version = TLSv1.2,TLSv1.3
```

Set ownership

```
chown -vR mysql:mysql /etc/mysql/ssl/
```

```
#### Restart and check for errors
```

```
systemctl restart mariadb journalctl -u mariadb
```

```
#### Setup on clients
```

from

copy /etc/mysql/ssl/ca-cert.pem

to client

```
cd /etc/mysql tar cvfz ssl.tar.gz ssl scp ssl.tar.gz 11trainingdo@ip:/tmp
```

```
sudo vi /etc/mysql/mariadb.conf.d/50-mysql-clients.cnf
```

Append/edit in [mysql] section:

MySQL Client Configuration

```
ssl-ca=/etc/mysql/ssl/ca-cert.pem
```

Force TLS version for client too

```
##tls_version = TLSv1.2,TLSv1.3
```

This option is disabled by default

```
ssl-verify-server-cert
```

only works if you have no self-signed certificate

```
ssl-verify-server-cert
```

```
### Test connection on client
```

```
mysql --ssl -uxyz -p -h mysql>status SSL: Cipher in use is TLS_AES_256_GCM_SHA384
```

```
### Force to use ssl
```

on server

now client can only connect, when using ssl

```
mysql> grant USAGE on . to remote@10.10.9.144 require ssl;
```

```
### On client to enable ssl by default for root
```

```
vi /root/.my.cnf [mysql] ssl
```

now mysql will always use ssl

```
mysql -uxyz -p -h10.10.9.110
```

```
### Ref
```

```
* https://www.cyberciti.biz/faq/how-to-setup-mariadb-ssl-and-secure-connections-from-clients/
```

Create User/Grant/Revoke - Management of users

Create user

```
create user training@localhost identified by 'deinpassword';
```

Drop user (=delete user)

```
drop user training@localhost
```

Change User (e.g. change authentication)

```
## change pass  
alter user training@localhost identified by 'newpassword';
```

Set global or db rights for a user

```
grant all on *.* to training@localhost  
## only a specific db  
grant all on mydb.* to training@localhost
```

Revoke global or db right from a user

```
revoke select on *.* from training@localhost  
## only from a specific db  
revoke select on training.* from training@localhost
```

Refs:

- <https://mariadb.com/kb/en/grant/#the-grant-option-privilege>
- <https://mariadb.com/kb/en/revoke/>

User- and Permission-concepts (best-practice)

Database - Objects

Create Database

Show all tables within db

show all tables in database

```
## connect with db training
mysql training
mysql> show tables;
|training|
```

describe

```
MariaDB [training]> describe mitarbeiter;
+-----+-----+-----+-----+-----+-----+
| Field | Type                | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+-----+
| id    | tinyint(3) unsigned | NO   | PRI | NULL    |       |
| name  | varchar(50)         | YES  |     | NULL    |       |
| vorname | varchar(30)        | YES  |     | NULL    |       |
+-----+-----+-----+-----+-----+-----+
3 rows in set (0.001 sec)
```

show create

```
MariaDB [training]> show create table mitarbeiter;
+-----+-----+
+-----+-----+
+-----+
| Table      | Create Table
+-----+-----+
+-----+-----+
+-----+-----+
| mitarbeiter | CREATE TABLE `mitarbeiter` (
  `id` tinyint(3) unsigned NOT NULL,
  `name` varchar(50) DEFAULT NULL,
  `vorname` varchar(30) DEFAULT NULL,
  PRIMARY KEY (`id`)
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 |
+-----+-----+
+-----+-----+
+-----+
1 row in set (0.000 sec)

<div class="page-break"></div>
```

```
## InnoDB - Storage Engine

### InnoDB - Storage Engine - Structure

![[InnoDB Structure](/images/InnoDB-Structure.jpg)]

<div class="page-break"></div>

### Important InnoDB - configuration - options to optimized performance

### Innodb buffer pool

* How much data fits into memory
* Free buffers = pages of 16 Kbytes
* Free buffer * 16Kbytes = free innodb buffer pool in KByte
```

pager grep -i 'free buffers' show engine innodb status \G Free buffers 7905 1 row in set (0.00 sec)

```
### Overview innodb server variables / settings

* https://dev.mysql.com/doc/refman/5.7/en/innodb-parameters.html

### Change innodb_buffer_pool
```

/etc/mysql/mysql.conf.d/mysqld.cnf

70-80% of memory on dedicated mysql

[mysqld] innodb-buffer-pool-size=6G

systemctl restart mysql

mysql mysql>show variables like 'innodb%buffer%';

```
### innodb_flush_method
```

Ideally O_DIRECT on Linux, but please test it, if it really works well.

```
### innodb_flush_log_at_trx_commit
```

When is flushing done from innodb_log_buffer to log. Default: 1 : After every commit -> best performance 2. -> once per second

Good to use 2, if you are willing to loose 1 second of data on powerfail

```
### innodb_flush_neighbors
```

on ssd disks set this to off, because there is no performance improvement

innodb_flush_neighbors=0

Default = 1

```
### skip-name-resolve.conf
```

work only with ip's - better for performance

/etc/my.cnf skip-name-resolve

```
* https://nixcp.com/skip-name-resolve/

### Ref:

* https://dev.mysql.com/doc/refman/5.7/en/innodb-buffer-pool-resize.html

### Privileges for show engine innodb status
```

show engine innodb status \G ERROR 1227 (42000): Access denied; you need (at least one of) the PROCESS privilege(s) for this operation

```
<div class="page-break"></div>

## Training Data

### Setup training data "contributions"

### Walkthrough

* Complete process takes about 10 minutes

```bash
cd /usr/src
apt update; apt install -y git
git clone https://github.com/jmetzger/dedupe-examples.git
cd dedupe-examples
cd mysql_example
Eventually you need to enter (in mysql_example/mysql.cnf)
```

```
Only necessary if you cannot connect to db by entering "mysql"
password=<your_root_pw>
./setup.sh
```

## Performance

### io-Last/CPU-Last

#### IO-gebundene - Last (Input/Output)

Gegeben wenn:

- Hoher waiting wert in top (wa-wert in CPU-Liste)
- + Hohelast 1,5, 15 min 1,2 1.5 2 (Load) -> top

#### CPU-Gebundene - Last

Gegeben wenn:

- NUR: Hohe Last -> Wert in top -> 2 1.5 0.5 (Load)
- Waiting-wert: 0

-

## Views and performance

### General

- SHOW CREATE VIEW
- Views can use 3 algorithms:
  - merge
  - simple rewrites (translates the query)
- temptable
  - Creates a temptable to retrieve information
  - In this case no indexes can be used
  - Shows up explain with : ``

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+ | id |
select_type | table | type | possible_keys | key | key_len | ref | rows | Extra | +-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+ | 1 | PRIMARY | | ALL | NULL | NULL |
NULL | NULL | 33 | NULL | | 2 | DERIVED | task | ALL | NULL | NULL | NULL | NULL | 33 | NULL | +-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

```
* undefined
* MySQL chooses, if to use merge or temptable
* prefers merge over temptable if possible
```

### Handling (best practice)

```
* You can define the algorithm when creating the view
* If you define merge and mysql cannot handle it
* you will get a warning
```

```
mysql> CREATE ALGORITHM=MERGE VIEW priority_counts AS SELECT priority_id, COUNT(1) AS quantity
FROM task GROUP BY priority_id; Query OK, 0 rows affected, 1 warning (0.12 sec)
```

```
mysql> SHOW WARNINGS; +-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+ | Level | Code | Message | +-----+-----+-----+-----+-----+
+-----+ | Warning | 1354 | View merge algorithm can't be used here for now
(assumed undefined algorithm) | +-----+-----+-----+-----+-----+
+-----+ 1 row in set (0.08 sec)
```

```
* Ref: https://dba.stackexchange.com/questions/54481/determining-what-algorithm-mysql-view-is-using
```

```
<div class="page-break"></div>
```

### Partitions and Explain

### Walkthrough

## EXPLAIN PARTITIONS

```
DROP TABLE IF EXISTS audit_log; CREATE TABLE audit_log (yr YEAR NOT NULL, msg VARCHAR(100) NOT NULL) ENGINE=InnoDB PARTITION BY RANGE (yr) (PARTITION p0 VALUES LESS THAN (2010), PARTITION p1 VALUES LESS THAN (2011), PARTITION p2 VALUES LESS THAN (2012), PARTITION p3 VALUES LESS THAN MAXVALUE); INSERT INTO audit_log(yr,msg) VALUES (2005,'2005'),(2006,'2006'),(2011,'2011'),(2020,'2020'); EXPLAIN PARTITIONS SELECT * from audit_log WHERE yr in (2011,2012)\G
```

```
Partitions sliced by hash of field
```

```
CREATE TABLE employees (id INT NOT NULL, fname VARCHAR(30), lname VARCHAR(30), hired DATE NOT NULL DEFAULT '1970-01-01', separated DATE NOT NULL DEFAULT '9999-12-31', job_code INT, store_id INT) PARTITION BY HASH(store_id) PARTITIONS 4;
```

```
<div class="page-break"></div>
```

```
3 Phases of DataSize
```

```
Phase 1: Table content is small (only some rows)
```

## table scan is quicker than index search

**e.g. 10 entries**

**so eventually index is not needed**

```
Phase 2: Index is good !!
```

## performance gain by using index

**Step 1: Obtaining id's from index (primary key id)**

**Step 2: Retrieving data**

```
Phase 3: Index is not improve performance / or would makes performance worse
```

Step 1: lookup in index: 1 70 1040 2100 35000 -> there is a lot of space (other rows) in between.

Step 2: Lookup data, but a lot lookups needed

-> random reads -> So mysql might be better off to do a table scan.

```
<div class="page-break"></div>
```

```
Backup and Restore (Point-In-Time aka PIT)
```

```
Backup with mysqldump - best practices
```

```
Useful options for PIT
```

**—quick not needed, because included in —opt which is enabled by default**

**on local systems using socket, there are no huge benefits concerning -compress**

**when you dump over the network use it for sure**

```
mysqldump --all-databases --single-transaction --gtid --master-data=2 --routines --events --flush-logs --compress > /usr/src/all-databases.sql;
```

```
With PIT_Recovery you can use --delete-master-logs
```

```
* All logs before flushing will be deleted
```

```
mysqldump --all-databases --single-transaction --gtid --master-data=2 --routines --events --flush-logs --compress --delete-master-logs > /usr/src/all-databases.sql;
```

```
Version with zipping
```

```
mysqldump --all-databases --single-transaction --gtid --master-data=2 --routines --events --flush-logs --compress | gzip > /usr/src/all-databases.sql.gz
```

```
Performance Test mysqldump (1.7 Million rows in contributions)
```

```
date; mysqldump --all-databases --single-transaction --gtid --master-data=2 --routines --events --flush-logs --compress > /usr/src/all-databases.sql; date Mi 20. Jan 09:40:44 CET 2021 Mi 20. Jan 09:41:55 CET 2021
```

```
Separated sql-structure files and data-txt files including master-data for a specific database
```

**backups needs to be writeable for mysql**

```
mkdir /backups chmod 777 /backups chown mysql:mysql /backups mysqldump --tab=/backups contributions mysqldump --tab=/backups --master-data=2 contributions mysqldump --tab=/backups --master-data=2 contributions > /backups/master-data.tx
```



```
<div class="page-break"></div>
```

```
Flashback
```

```
* Redoes insert/update/delete entries from binlog (binlog_format = 'ROW')
```

```
Referenz:
```

```
* https://mariadb.com/kb/en/flashback/
```

```
<div class="page-break"></div>
```

```
mariabackup
```

```
Walkthrough
```

## user eintrag in /root/.my.cnf

```
[mariabackup] user=root
```

**pass is not needed here, because we have the user root with  
unix\_socket - auth**

```
mkdir /backups
```

## target-dir needs to be empty or not present

```
mariabackup --target-dir=/backups/20210120 --backup
```

## apply ib\_logfile0 to tablespaces

## after that ib\_logfile0 -> 0 bytes

```
mariabackup --target-dir=/backups/20210120 --prepare
```

## Recover

```
systemctl stop mariadb mv /var/lib/mysql /var/lib/mysql.bkup mariabackup --target-dir=/backups/20200120
--copy-back chmod -R mysql:mysql /var/lib/mysql systemctl start mariadb
```

```
Ref.
```

```
https://mariadb.com/kb/en/full-backup-and-restore-with-mariabackup/
```

```
<div class="page-break"></div>
```

```
Use xtrabackup for MariaDB 5.5
```

```
For mariadb 5.5 you can use xtrabackup instead of mariabackup
```

```
* https://www.percona.com/doc/percona-xtrabackup/2.4/index.html
```

```
<div class="page-break"></div>
```

```
Optimal use of indexes
```

```
Index and Functions (Cool new feature in MySQL 5.7)
```

```
No index can be used on an index:
```

```
explain select * from actor where upper(last_name) like 'A%'; +----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+-----+-----+-----+ | id | select_type | table |
partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra | +----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+ | 1 | SIMPLE |
actor | NULL | ALL | NULL | NULL | NULL | NULL | 200 | 100.00 | Using where | +----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
```

```
Workaround with virtual columns (possible since mysql 5.7)
```

## 1. Create Virtual Column with upper

```
alter table sakila add idx_last_name_upper varchar(45) GENERATED ALWAYS AS upper(last_name);
```

## 2. Create an index on that column

```
create index idx_last_name_upper on actor (last_name_upper);
```

```
Now we try to search the very same
```

```
explain select * from actor where last_name_upper like 'A%'; +----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+-----+-----+-----+ | id |
select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra | +----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+ | 1 | SIMPLE | actor | NULL | range | idx_last_name_upper | idx_last_name_upper |
183 | NULL | 7 | 100.00 | Using where | +----+-----+-----+-----+-----+-----+
---+-----+-----+-----+-----+-----+-----+-----+ 1 row in set, 1 warning (0.00
sec)
```

```
Preview MySQL 8
```

```
* MySQL 8 support functional indexes
```

```
<div class="page-break"></div>
```

```
Index and Likes
```

```
1. like 'Will%' - Index works
```

```
explain select last_name from donors where last_name like 'Will%';
```

```
2. like '%iams' - Index does not work
```

```
-- because like starts with a wildcard explain select last_name from donors where last_name like '%iams';
```

```
3. How to fix 3, if you are using this often ?
```

## Walkthrough

### Step 1: modify table

```
alter table donors add last_name_reversed varchar(70) GENERATED ALWAYS AS (reverse(last_name));
create index idx_last_name_reversed on donors (last_name_reversed);
```

### besser - Variante 2 - untested

```
alter table donors add last_name_reversed varchar(70) GENERATED ALWAYS AS (reverse(last_name)), add
index idx_last_name_reversed on donors (last_name_reversed);
```

### Step 2: update table - this take a while

```
update donors set last_name_reversed = reversed(last_name)
```

### Step 3: work with it

```
select last_name,last_name_reversed from donor where last_name_reversed like reverse('%iams');
```

## Version 2 with pt-online-schema-change

```
<div class="page-break"></div>
```

```
profiling-get-time-for-execution-of.query
```

```
* Get better values, how long queries take
```

```
Example
```

```
set profiling = 1
```

## Step 2 - Execute query

```
select last_name as gross from donors where last_name like lower('WILLI%')
```

## Step 3 - Show profiles

```
show profiles; +-----+-----+-----+
+-----+ | Query_ID | Duration | Query | +-----+-----+-----+
+-----+ | 1 | 0.01993525 | select last_name as gross from
donors where last_name like lower('WILLI%') | 4 rows in set, 1 warning (0.00 sec)
```

## Step 4 - Show profile for a specific query

```
mysql> show profile for query 1; +-----+-----+ | Status | Duration | +-----+
+-----+-----+ | starting | 0.000062 | | checking permissions | 0.000006 | | Opening tables | 0.000021 | |
init | 0.000017 | | System lock | 0.000007 | | optimizing | 0.000007 | | statistics | 0.000083 | | preparing |
0.000012 | | executing | 0.000004 | | Sending data | 0.022251 | | end | 0.000005 | | query end | 0.000008 | |
closing tables | 0.000007 | | freeing items | 0.001792 | | cleaning up | 0.000016 | +-----+-----+
+-----+ 15 rows in set, 1 warning (0.00 sec)
```

```
<div class="page-break"></div>
```

```
Find out cardinality without index
```

```
Find out cardinality without creating index
```

```
select count(distinct donor_id) from contributions;
```

```
select count(distinct(vendor_city)) from contributions; +-----+
count(distinct(vendor_city)) | +-----+ | 1772 | +-----+ 1
row in set (4.97 sec)
```

```
<div class="page-break"></div>
```

```
Monitoring
```

```
What to monitor?
```

```
What to monitor
```

```
System
```

- \* Last auf dem System (top)
- \* Festplatte (z.B. 85% voll ?) df /var/lib/mysql
- \* Swap (Wenn geswappt wird ist Hopfen und Malz verloren)

```
Erreichbarkeit
```

```
* Server per ping erreichen (mysqladmin ping -h ziel-ip)
* Einlogbar ? (myadmin ping -h ziel-ip -u control_user
```

```
Platte aka IO-Subsystem (iostats)
```

```
* http://schulung.t3isp.de/documents/pdfs/mysql/mysql-performance.pdf
```

```
| -- | -- | -- |
| ----- |:-----:| ----:|
| Read/Write requests | IOPS (Input/Output operations per second) | -- |
| Average IO wait | Time that queue operations have to wait for disk access | --
|
| Average Read/Write time | Time it takes to finish disk access operations (latency) | |
|---|---|---|
| Read/Write bandwidth | Data transfer from and towards your disk | -- |
```

```
Gneral mysql metrics
```

```
mysql -E -e "select variable_value from information_schema.session_status where variable_name =
'uptime'";
```

## max connections

```
MariaDB [(none)]> show status like 'max_used_connections'; +-----+-----+ |
Variable_name | Value | +-----+-----+ | Max_used_connections | 1 | +-----+
--+-----+ 1 row in set (0.001 sec)
```

```
MariaDB [(none)]> show variables like 'max_connections'; +-----+-----+ | Variable_name |
Value | +-----+-----+ | max_connections | 151 | +-----+-----+ 1 row in set
(0.001 sec)
```

mysqladmin status

## you will find uptime here in seconds

```
| Metric | Comments | Suggested Alert |
| ----- |:-----:| ----:|
| Uptime | Seconds since the server was started. We can use this to detect
respawns. | When uptime is < 180. (seconds) |
| Threads_connected | Number of clients currently connected. If none or too high,
something is wrong. | None |
| Max_used_connections | Max number of connections at a time since server started.
(max_used_connections / max_connections) indicates if you could run out soon of
connection slots. | When connections usage is > 85%. |
| Aborted_connects | Number of failed connection attempts. When growing over a
period of time either some credentials are wrong or we are being attacked. | When
aborted connects/min > 3. |
```

#### #### InnoDB

Metric	Comments	Suggested Alert
Innodb_row_lock_waits	Number of times InnoDB had to wait before locking a row.	None
Innodb_buffer_pool_wait_free	Number of times InnoDB had to wait for memory pages to be flushed. If too high, innodb_buffer_pool_size is too small for current write load.	None

#### #### Query tracking

Metric	Comments	Suggested Alert
Slow_queries	Number of queries that took more than long_query_time seconds to execute. Slow queries generate excessive disk reads, memory and CPU usage. Check slow_query_log to find them.	None
Select_full_join	Number of full joins needed to answer queries. If too high, improve your indexing or database schema.	None
Created_tmp_disk_tables	Number of temporary tables (typically for joins) stored on slow spinning disks, instead of faster RAM.	None
(Full table scans) Handler_read%	Number of times the system reads the first row of a table index. (if 0 a table scan is done - because no key was read). Sequential reads might indicate a faulty index.	None

#### #### Track Errors

journalctl -u mariadb | grep -i Error

#### #### Ref

\* <https://blog.serverdensity.com/how-to-monitor-mysql/>

#### #### Monitoring with pmm (Percona Management Monitoring)

<https://pmmdemo.percona.com>

[Documentation] (<https://www.percona.com/doc/percona-monitoring-and-management/2.x/details/commands/pmm-admin.html>)

<div class="page-break"></div>

#### ## Replication

#### ### Slave einrichten -gtid

#### ### Step 1: mariabackup on master

mkdir /backups

## target-dir needs to be empty or not present

```
mariabackup --target-dir=/backups/20210121 --backup
```

## apply ib\_logfile0 to tablespaces

## after that ib\_logfile0 -> 0 bytes

```
mariabackup --target-dir=/backups/20210121 --prepare
```

```
Step 2: Transfer to new slave (from master)
```

## root@master:

```
rsync -e ssh -avP /backups/mysqldumpdir/20210121 kurs@10.10.9.144:/home/kurs/
```

```
Step 3: Setup replication user on master
```

## as root@master

```
##mysql> CREATE USER repl@'10.10.9.%' IDENTIFIED BY 'password'; GRANT REPLICATION SLAVE ON . TO 'repl'@'10'
```

```
Step 3a (Optional): Test repl user (connect) from slave
```

## as root@slave

## you be able to connect to

```
mysql -urepl -p -h10.10.9.110
```

## test if grants are o.k.

```
show grants
```

```
Step 4a: Set server-id on master -> 1
```

```
[mysqld] server-id=1
```

```
systemctl restart mariadb
```

```
Step 4b: Set server-id on slave -> 3 + same config as server 1
```

```
[mysqld] server-id = 3
```

## activate master bin log, if this slave might be a master later

```
log_bin = /var/log/mysql/mysql-bin.log
```

```
systemctl restart mariadb
```

## auf dem master config mit rsync rüberschrieben

**root@master**

```
rsync -e ssh -avP /etc/mysql/mariadb.conf.d/z_uniruhr.cnf kurs@10.10.9.144:/home/kurs/
```

**root@slave**

```
mv /home/kurs/z_uniruhr.cnf /etc/mysql/mariadb.conf.d/ chown root:root /etc/mysql/mariadb.conf.d
systemctl restart mariadb
```

```
Step 5: Restore Data on slave
```

```
systemctl stop mariadb mv /var/lib/mysql /var/lib/mysql.bkup4 mariabackup --target-dir=/backups/20210121
--copy-back chown -R mysql:mysql/var/lib/mysql systemctl start mariadb
```

```
Step 6: master.txt for change command
```

**root@slave**

```
$ cat xtrabackup_binlog_info mariadb-bin.000096 568 0-1-2
```

```
SET GLOBAL gtid_slave_pos = "0-1-2";
```

## /root/master.txt

## get information from master-databases.sql dump

```
CHANGE MASTER TO MASTER_HOST="10.10.9.110", MASTER_PORT=3306, MASTER_USER="repl",
MASTER_PASSWORD="password", MASTER_USE_GTID=slave_pos;
```

```
mysql < master.txt
```

## or: copy paste into mysql>

**mysql>**

```
start slave
```

## in mysql -> show slave status

```
mysql>show slave status
```

## Looking for



Slave\_IO\_Running: Yes Slave\_SQL\_Running: Yes

```
Walkthrough
```

```
https://mariadb.com/kb/en/setting-up-a-replication-slave-with-mariabackup/
```

```
<div class="page-break"></div>
```

```
Slave einrichten - master_pos
```

```
Step 1: mysqldump on master
```

```
mkdir -p /backups/mysqldumpdir
```

**in version 5.5. there is not --git so use it without --gtid**

```
mysqldump --all-databases --single-transaction --master-data=2 --routines --events --compress > /backups/mysqldumpdir/master-databases.sql;
```

```
Step 2: Transfer to new slave (from master)
```

**root@master:**

```
rsync -e ssh -avP /backups/mysqldumpdir/master-databases.sql kurs@10.10.9.144:/home/kurs/
```

```
Step 3 (Optional): Be sure that slave is really fresh (no data yet)
```

**if old not wanted data is present, e.g. other databases, start with fresh-installation by so:**

**as root**

```
cd /var/lib mv mysql mysql.bkup mariadb-install-db --user=mysql
```

```
Step 4: Setup replication user on master
```

**as root@master**

```
##mysql> CREATE USER repl@'10.10.9.%' IDENTIFIED BY 'password'; GRANT REPLICATION SLAVE ON . TO 'repl'@'10
```

```
Step 4a (Optional): Test repl user (connect) from slave
```

**as root@slave**

**you be able to connect to**

mysql -urepl -p -h10.10.9.110

**test if grants are o.k.**

show grants

```
Step 5a: Set server-id on master -> 1
```

[mysqld] server-id=1

systemctl restart mariadb

```
Step 5b: Set server-id on slave -> 2 + same config as server 1
```

[mysqld] server-id = 2

**activate master bin log, if this slave might be a master later**

log\_bin = /var/log/mysql/mysql-bin.log

systemctl restart mariadb

**auf dem master config mit rsync rüberschrieben**

**root@master**

rsync -e ssh -avP /etc/mysql/mariadb.conf.d/z\_uniruhr.cnf [kurs@10.10.9.144](mailto:kurs@10.10.9.144):/home/kurs/

**root@slave**

mv /home/kurs/z\_uniruhr.cnf /etc/mysql/mariadb.conf.d/ chown root:root /etc/mysql/mariadb.conf.d

systemctl restart mariadb

```
Step 6: Restore Data on slave
```

**root@slave**

cd /home/kurs mysql < master-databases.sql

```
Step 7: master.txt for change command
```

**root@slave**

**/root/master.txt**

## **get information from master-databases.sql dump**

```
CHANGE MASTER TO MASTER_HOST="10.10.9.110", MASTER_PORT=3310, MASTER_USER="repl",
MASTER_PASSWORD="password", MASTER_LOG_FILE='mysqld-bin.000001', MASTER_LOG_POS=568;
```

## **Version 1**

```
mysql < master.txt
```

**or: copy paste into mysql>**

**in mysql -> show slave status**

```
mysql>show slave status
```

## **Looking for**

Slave\_IO\_Running: Yes Slave\_SQL\_Running: Yes

```
Step 8: not working on 5.5.
```

Switch to using gtid later on:

```
show slave status; # look for using_gtid stop slave; CHANGE MASTER TO MASTER_USE_GTID = slave_pos;
show slave status; # look for using_gtid start slave;
```

```
Walkthrough
```

```
https://mariadb.com/kb/en/setting-up-a-replication-slave-with-mariabackup/
```

```
<div class="page-break"></div>
```

```
MaxScale installieren
```

```
Why do Loadbalancing with MaxScale ?
```

- \* Cluster node transparent to application
  - \* Application does not see single nodes
- \* If one node fails you will have no downtime
  - \* In opposite: To talking to this node directly

```
License Implications since 2.x
```

- \* MariaDB MaxScale >= 2.0 is licensed under MariaDB BSL.

```
* maximum of three servers in a commercial context.
 * Any more, and you'll need to buy their commercial license.

* MariaDB MaxScale 2.1.0 will be released under BSL 1.1 from the start

* Each release transitions in about max 4 years to GPL
```

### ### The MaxScale load-balancer and its components

```
* Routers
* Listeners
* Filters
* Servers (backend database server)
```

#### #### Filters

```
* Logging Filters
* Statement rewriting filters
* Result set manipulation filters
* Firewall filter
* Pipeline control filters
 * e.g. tee and send to a second server

* Ref: https://mariadb.com/kb/en/mariadb-maxscale-25-regex-filter/
```

### ### Documentation - maxctrl

```
* https://mariadb.com/kb/en/mariadb-maxscale-25-maxctrl/
```

### ### Installation and Setup

#### #### Installation

apt update apt install apt-transport-https curl

## Setting up the repos

curl -sS [https://downloads.mariadb.com/MariaDB/mariadb\\_repo\\_setup](https://downloads.mariadb.com/MariaDB/mariadb_repo_setup) | sudo bash

## Installing maxscale

apt install maxscale

```
Setup (Part 1: MaxScale db-user)

* Do this on one of the galera nodes
* Adjust IP !!

```bash
```

```

## IP FROM MAXSCALE
## Setup privileges on cluster nodes
## It is sufficient to set it on one node, because
## it will be synced to all the other nodes
## on node 1
CREATE USER 'maxscale'@'10.10.11.139' IDENTIFIED BY 'P@ssw0rd';
##
GRANT SELECT ON mysql.db TO 'maxscale'@'10.10.11.139';
GRANT SELECT ON mysql.user TO 'maxscale'@'10.10.11.139';
GRANT SELECT ON mysql.tables_priv TO 'maxscale'@'10.10.11.139';
##
GRANT SELECT ON mysql.columns_priv TO 'maxscale'@'10.10.11.139';
GRANT SELECT ON mysql.proxies_priv TO 'maxscale'@'10.10.11.139';
##
GRANT SHOW DATABASES ON *.* TO 'maxscale'@'10.10.11.139';
## Needed for maxscale
GRANT SELECT ON mysql.procs_priv TO 'maxscale'@'10.10.11.139';
GRANT SELECT ON mysql.roles_mapping TO 'maxscale'@'10.10.11.139';

## Additionally for cluster operations (rejoin, switchover, failover for master/slave
replications
## these permissions are needed
GRANT super, reload, process, show databases, event on *.* to
'maxscale'@'10.10.11.139';
## GRANT select on mysql.user to 'maxscale'@'10.10.11.139';

```

```

## On maxscale - server
apt update
apt install mariadb-client
## Test the connection
## Verbindung sollte aufgebaut werden
mysql -u maxscale -p -h <ip-eines-der-nodes>
mysql>show databases

```

SETUP (PART 2: CONFIGURATION)

```

## /etc/maxscale.cnf

[maxscale]

threads=auto
syslog=0
maxlog=1
log_warning=1
log_notice=1
log_info=0
log_debug=0

[TheMonitor]
type=monitor
module=mariadbmon

```

```
servers=server1,server2,server3
user=maxscale
password=P@ssw0rd
auto_rejoin=true
auto_failover=true
```

```
[RW-Split-Router]
type=service
router=readwritesplit
servers=server1,server2,server3
user=maxscale
password=P@ssw0rd
max_slave_connections=100%
```

```
[RW-Split-Listener]
type=listener
service=RW-Split-Router
protocol=MariaDBClient
port=3306
```

```
[server1]
type=server
address=142.93.98.60
port=3306
protocol=MariaDBBackend
```

```
[server2]
type=server
address=142.93.103.153
port=3306
protocol=MariaDBBackend
```

```
[server3]
type=server
address=142.93.103.246
port=3306
protocol=MariaDBBackend
```

```
## Start
```

```
systemctl start maxscale
```

```
## What does the log say ?
```

```
## /var/log/maxscale/maxscale.log
```

maxctrl

```
maxctrl list servers
maxctrl show server server1
```

```
maxctrl list services  
maxctrl show service ReadWrite-Split-Router
```

Reference: MaxScale-Proxy mit Monitoring

[MaxScale MariaDB-Monitor](#)

Walkthrough:Automatic Failover Master Slave

<https://mariadb.com/kb/en/mariadb-maxscale-25-automatic-failover-with-mariadb-monitor/>

Tools

Percona-toolkit-Installation

Walkthrough

```
## Howto
## https://www.percona.com/doc/percona-toolkit/LATEST/installation.html

## Step 1: repo installieren mit deb -paket
wget https://repo.percona.com/apt/percona-release_latest.focal_all.deb;
apt update;
apt install -y curl;
dpkg -i percona-release_latest.focal_all.deb;
apt update;
apt install -y percona-toolkit;
```

pt-query-digest - analyze slow logs

Requires

- Install percona-toolkit

Usage

```
## first enable slow_query_log
set global slow_query_log = on
set global long_query_time = 0.2
## to avoid, that i have to reconnect with new session
set session long_query_time = 0.2

## produce slow query - for testing
select * from contributions where vendor_last_name like 'W%';
mysql > quit

##
cd /var/lib/mysql
## look for awhile with -slow.log - suffix
pt-query-digest mysql-slow.log > /usr/src/report-slow.txt
less report-slow.txt
```

pt-online-schema-change howto

Requirements

- Install percona-toolkit

What does it do ?

```
## Altering table without blocking them
## Do a dry-run beforehand
pt-online-schema-change --alter "ADD INDEX idx_city (city)" --dry-run
D=contributions,t=donors
##
pt-online-schema-change --alter "ADD INDEX idx_city (city)" --execute
D=contributions,t=donors
```

Problems -> high cpu load

```
## fine - tune params
## e.g. --max-load
## refer to docs
https://www.percona.com/doc/percona-toolkit/3.0/pt-online-schema-
change.html#:~:text=pt%2Donline%2Dschema%2Dchange%20works%20by%20creating%20an%20empty,i
```

Extras

User Variables

```
## only valid within one session
set @host='localhost';

## You can use it in select
select @host;

## You can use it in the where clause
select mysql.user where host=@host;

## not possible to use it within create user
## DOES NOT WORK !
set @mypass='password';
create user someuser@somehost identified by @mypass;
```

Installation sakila-db

```
cd /usr/src
wget https://downloads.mysql.com/docs/sakila-db.tar.gz
tar xzvf sakila-db.tar.gz

cd sakila-db
mysql < sakila-schema.sql
mysql < sakila-data.sql
```

Documentation

Server System Variables

- https://mariadb.com/kb/en/server-system-variables/#bind_address

MySQL - Performance - PDF

- <http://schulung.t3isp.de/documents/pdfs/mysql/mysql-performance.pdf>

Source-Code MariaDB

- <https://github.com/MariaDB/server>

Diagnosis and measurement of performance

Best practices to narrow down performance problems

Pre-Requisites

- System is slow

Analyze - Checklist - Step 1

```
## Are there slow queries ?
## look for time
show full processlist

### or time - in seconds
select * from information_schema.processlist where time > 10;
```

Re-Execute SELECT or where from UPDATE / DELETE

```
## Is it still slow ?
## Eventually kill
mysql>show processlist
mysql>--kill <Thread-id>
mysql>-- example
mysql>kill 44
```

Explain what is going on

```
Explain Select....
```

Performance and optimization of SQL statements

Do not use '*' whenever possible

Why ?

- You are adding .. to the server:
 - I/O
 - memory
 - CPU
- You are preventing covering indexes

Walkthrough. (Look at the time)

Using '*'

```
## using '*'
pager grep "rows in set";
select * from donors where last_name like 'Willia%'; select * from donors where
last_name like 'Willia%';
-- time between 0.02 and 0.04 secs
-- 2424 rows in set (0.02 sec)
-- reset pager
pager

## corresponding Explain (QEP)
explain select * from donors where last_name like 'Willia%';
+---+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key |
| key_len | ref | rows | filtered | Extra |
+---+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | donors | NULL | range | donors_donor_info |
donors_donor_info | 213 | NULL | 4748 | 100.00 | Using index condition |
+---+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.00 sec)
```

using specific fields

```
pager grep 'rows in set'; select last_name,first_name from donors where last_name like
'Willia%'; pager;
PAGER set to 'grep 'rows in set''
2424 rows in set (0.01 sec)
```

```
explain select last_name,first_name from donors where last_name like 'Willia%';
+---+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key |
| key_len | ref | rows | filtered | Extra |
+---+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.00 sec)
```



```

+-----+-----+-----+-----+-----+
| 1 | SIMPLE | donors | NULL | range | donors_donor_info |
donors_donor_info | 213 | NULL | 4748 | 100.00 | Using where; Using index |
+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.00 sec)

```

- Uses cover index (indicator in Extra: using index)

Ref:

- <https://www.oreilly.com/library/view/high-performance-mysql/9780596101718/ch04.html>

Be aware of subselects - Example 1

Optimizer-hints (and why you should not use them)

Tell the optimizer what to do and what not to do

- <https://dev.mysql.com/doc/refman/5.7/en/optimizer-hints.html#optimizer-hints-syntax>

Replication

Replikation Read/Write

- <https://proxysql.com/blog/configure-read-write-split/>

Performance

Best Practices

Indexes

2 Indexes vs. Combined Index

- In most cases a combined index is better than 2 indexes.

Joins

Field-Type

- Do not use varchar() or char() aka string types of join field
- better: integer (unsigned) && same size
 - e.g. actor_id id int unsigned

Views

General

- Only use views with merge
- NO temptable please, these CANNOT be indexed.

Where

No functions in where please

- Why ? Index cannot be used.
- example:
 - select first_name from actor where upper(first_name) like 'A%'

Alternative solution

- use a virtual field and index virtual field (possible from mysql > 5.7)
- Massive improvements in mysql 8

Example sys-schema and Reference

Examples

```
mysql> select * from sys.host_summary\G
***** 1. row *****
      host: localhost
    statements: 1347
 statement_latency: 7.55 m
statement_avg_latency: 336.50 ms
      table_scans: 15
       file_ios: 612857
  file_io_latency: 1.66 m
current_connections: 1
total_connections: 7
      unique_users: 1
      current_memory: 0 bytes
total_memory_allocated: 0 bytes
1 row in set (0.01 sec)
```

Ref:

- <https://github.com/mysql/mysql-sys/blob/master/README.md>

Change schema online (pt-online-schema-change)

- <https://www.percona.com/doc/percona-toolkit/3.0/pt-online-schema-change.html>

Optimizer-Hints

Tell the optimizer what to do and what not to do

- <https://dev.mysql.com/doc/refman/5.7/en/optimizer-hints.html#optimizer-hints-syntax>

Documentation / Literature

Effective MySQL

- <https://www.amazon.com/Effective-MySQL-Optimizing-Statements-Oracle/dp/0071782796>

Last Training

- <https://github.com/jmetzger/training-mysql-developers-basics>

MySQL - Performance - PDF

- <http://schulung.t3isp.de/documents/pdfs/mysql/mysql-performance.pdf>

MariaDB Galera Cluster

- <http://schulung.t3isp.de/documents/pdfs/mariadb/mariadb-galera-cluster.pdf>

MySQL Galera Cluster

- <https://galeracluster.com/downloads/>

Questions and Answers

Questions and Answers

1. Do you recommend Aurora

```
In my current humble opinion Aurora is a double edged sword.  
Aurora looks promising for scalability, but a lot of stuff is modified  
mysql-stuff and in my opinion has a lot of restrictions.
```

```
You should be aware, that moving to Aurora might be a task  
and reverting back even more.
```

- Refer to: <https://ahmedahamid.com/aurora-mysql/>

I would like to point you to a performance measurement report here:

- <https://galeracluster.com/2019/09/everdata-reports-galera-cluster-outshines-amazon-aurora-and-rds/>

2. Get rid of unattended - upgrades problem (dirty hack)

```
ps aux | grep unatt  
kill <process-id-von-unattended-upgrades>
```

3. Archive Data

```
https://www.percona.com/doc/percona-toolkit/LATEST/pt-archiver.html
```

4. Does innodb do defragmentation by itself ?

```
## Some background while doing research.  
## Nil performance benefits of defragmentation in index.
```

```
https://stackoverflow.com/questions/48569979/mariadb-table-defragmentation-using-optimize
```

5. Defragmentation

```
## Optimize table
ALTER TABLE contributions engine = InnoDB

## mariadb has a patch for defragmentation
https://mariadb.org/defragmenting-unused-space-on-innodb-tablespace/

## alter table xyz engine=InnoDB - defragmentations
## but is also invasive.
## with ibdata1 innodb_file_per_table it lets the size grow
```

6. Is it possible to do select, update, deletes without using innodb_buffer in specific

```
No, this is not possible
```

7. Unit test framework in MySQL

```
No, there is no testing framework with MySQL
```

8. MariaDB - Advantages

- flashback
- Verschlüsselung von Tabellen // mariabackup
- Einige Storage Engine (Aria -> MyISAM - crash-recovery)
- JSON anders implementiert
- galera
- feature: defragmentation

```
MySQL 8 does not:
decode
set profiling (still available but deprecated )
```

9. Select without locking

```
SET TRANSACTION ISOLATION LEVEL READ UNCOMMITTED ;
BEGIN ;
SELECT * FROM TABLE_NAME ;
COMMIT ;
```

migration-mysql-update-5.6->5.7

=====

1. Sicherung.

xtrabackup

Mysqldump

16 GB

———

1.

Neue Location -> 5.6

<- Xtrbackup

Server runterfahren

Update 5.7

Fahrt den Server wieder hoch

2. Source-Host (Old Host) -> mysqldump

Neuen -> Installation von MySQL 5.7

Test-einspielen.

< mysqldump

4-5 Stunden.

-> Konfiguration von mysql -> was wollt ihr übernehmen.

3. Replications - Slave auf neuem System -> 5.7

Hängt in den Master.

Sicheren Transport

-> ssh -tunnel .

-> Firewall-Regeln.

-> ssl -absicherung

MySQL Do-Nots

mysql-do-nots

1. No function in where (column_name)

```
## Never use a function for the column name in where
## e.g.
select * from donors where upper(last_name) like 'Willia%'
```

Why ?

- Not index can be used

```
## Not filtering possible by indx -> possible_keys -> NULL
explain select last_name from donors where upper(last_name) like 'WILLI%';
+----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key |
key_len | ref | rows | filtered | Extra |
+----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | donors | NULL | index | NULL | donors_donor_info |
687 | NULL | 701948 | 100.00 | Using where; Using index |
+----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.00 sec)
```